

# Wine Quality Prediction Project

## Objective

- The objective of this project is to predict the quality of wine as bad-good (scale 1-10) using machine learning algorithms in Python. The dataset is used to create models to predict the quality of wine through different parameters like fixed acidity, volatile acidity etc.
- Multiple Classification models are applied and the accuracy score for the different models are compared.
- The model with best accuracy score will be used to help predict the quality of the wine as Good or Bad

In [1]:



```
1 # importing the needed libraries
2 import numpy as np
3 import pandas as pd
4
5 # for visualization
6 import matplotlib.pyplot as plt
7 import seaborn as sns
8 %matplotlib inline
9
10 # for outlier detection and removal
11 from scipy.stats import zscore
12
13 # for detection of multicollinearity using VIF
14 from statsmodels.stats.outliers_influence import variance_inflation_factor
15 from statsmodels.tools.tools import add_constant
16
17 # for data scaling
18 from sklearn.preprocessing import StandardScaler
19
20 # for model building and evaluation
21 from sklearn.model_selection import train_test_split
22 from sklearn.tree import DecisionTreeClassifier
23 from sklearn.ensemble import RandomForestClassifier
24 from sklearn.linear_model import LogisticRegression
25 from sklearn.neighbors import KNeighborsClassifier
26 from sklearn.model_selection import cross_val_score
27 from sklearn.metrics import confusion_matrix, classification_report
28 from sklearn.metrics import roc_curve, auc
29 from sklearn import metrics
30
31 # for hyper parameter tuning
32 import time
33 from sklearn.model_selection import GridSearchCV
34
35 # for display of model comparisons
36 from IPython.display import HTML
37
38 # ignoring warnings
39 import warnings
40 warnings.filterwarnings('ignore')
```

## Reading the Data

In [2]:

```
1 # reading the wine quality dataset
2 wine_data = pd.read_csv("QualityPrediction.csv")
```

## Exploratory Data Analysis

In [3]:

```
1 # inspecting the first 5 rows of wine quality data
2 wine_data.head()
```

Out[3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcoh
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9

In [4]:

```
1 # displaying all the features(columns) in the wine_data
2 print(wine_data.columns)
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'densit
y',
      'pH', 'sulphates', 'alcohol', 'quality'],
      dtype='object')
```

In [5]:

```
1 # a peek into the dataset
2 wine_data.shape
3
```

Out[5]:

(1599, 12)

The dataset has 1599 observations and 12 columns

In [6]:



```
1 # Looking at the unique values of the target variable
2 print("The unique values of the target variable (Quality) are: "+ str(wine_data['quality']
```

The unique values of the target variable (Quality) are: [5 6 7 4 8 3]

## Checking for Null Values

In [7]:



```
1 # identifying the data types of the columns and the number of missing values in each column
2 display(pd.DataFrame({'DataType':wine_data.dtypes,'Missing Values':wine_data.isnull().sum()})
```

	DataType	Missing Values
fixed acidity	float64	0
volatile acidity	float64	0
citric acid	float64	0
residual sugar	float64	0
chlorides	float64	0
free sulfur dioxide	float64	0
total sulfur dioxide	float64	0
density	float64	0
pH	float64	0
sulphates	float64	0
alcohol	float64	0
quality	int64	0

All the columns are of numeric type and there are no missing values

In [8]:



```
1 ## Describing the data -Five point summary of the data
2 wine_data.describe()
```

Out[8]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467000
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895700
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000

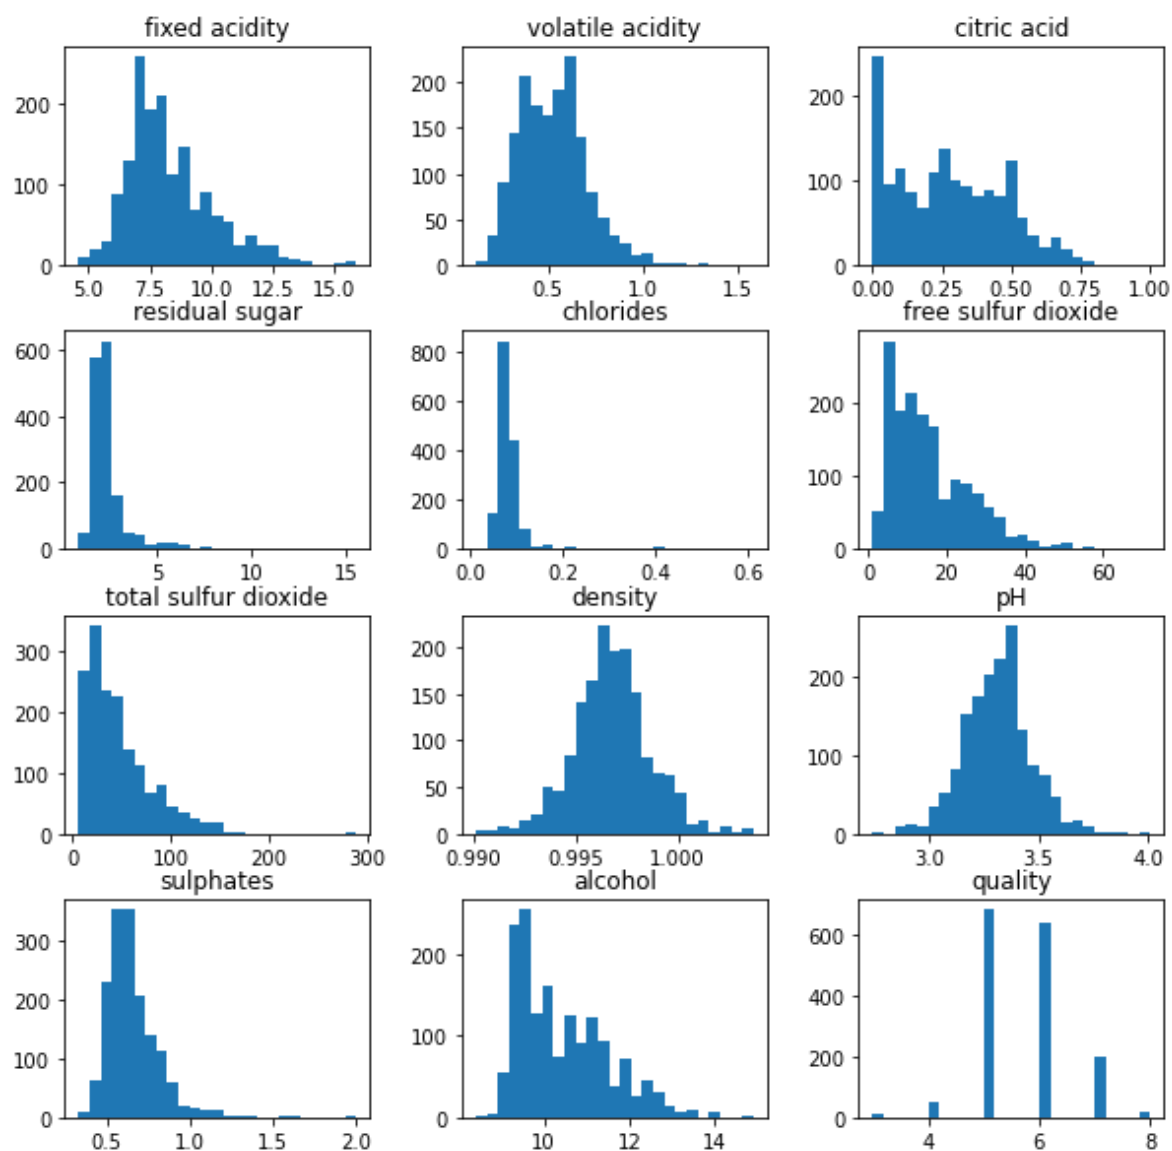


### Visualizing the data distribution

In [9]:



```
1 # Histogram Distribution of the data
2 wine_data.hist(bins=25,figsize=(10,10),grid=False)
3 # display histogram
4 plt.show()
```



The histograms show the distribution of each feature. The histogram also gives us an idea about the outliers in the given data set.

## Plotting the Relationship between Independent Variables and the Quality(Target Variable)


### a) Fixed Acidity and Quality

In [10]:

```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'fixed acidity', data = wine_data)
3 '''The relationship between fixed acidity and quality is ambiguous'''
```

Out[10]:

'The relationship between fixed acidity and quality is ambiguous'



### b) Volatile Acidity and Quality


In [11]:



```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'volatile acidity', data = wine_data)
3 '''As the quality of wine increases, the volatile acidity decreases'''
```

Out[11]:

'As the quality of wine increases, the volatile acidity decreases'



### c) Citric acid and Quality


In [12]:



```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'citric acid', data = wine_data)
3 '''As the quality of wine increases, the citric acid content increases'''
```

Out[12]:

'As the quality of wine increases, the citric acid content increases'



#### d) Residual Sugar and Quality




In [13]:



```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'residual sugar', data = wine_data)
3 '''The is no significant effect of residual sugar on the quality of wine'''
```

Out[13]:

'The is no significant effect of residual sugar on the quality of wine'



### e) Chlorides and Quality


In [14]:



```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'chlorides', data = wine_data)
3 '''As the quality of wine increases, the chloride content decreases'''
```

Out[14]:

'As the quality of wine increases, the chloride content decreases'




#### f) Free Sulfur Dioxide and Quality

In [15]:

```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'free sulfur dioxide', data = wine_data)
3
```

Out[15]:

&lt;AxesSubplot:xlabel='quality', ylabel='free sulfur dioxide'&gt;




### g) Total Sulfur Dioxide and Quality

In [16]:

```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'total sulfur dioxide', data = wine_data)
3
```

Out[16]:

&lt;AxesSubplot:xlabel='quality', ylabel='total sulfur dioxide'&gt;



## h) Density and Quality


In [17]:



```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'density', data = wine_data)
3 '''Density is statistically same irrespective of the quality of wine'''
```

Out[17]:

'Density is statistically same irrespective of the quality of wine'




### i) pH and Quality

In [18]:

```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'pH', data = wine_data)
3 '''There is a very slight decrease in the pH of the wine as quality increases'''
```

Out[18]:

'There is a very slight decrease in the pH of the wine as quality increases'




### k) Sulphates and Quality

In [19]:

```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'sulphates', data = wine_data)
3 '''There is an increase in the sulphates content as the quality of wine increases'''
```

Out[19]:

'There is an increase in the sulphates content as the quality of wine increases'




## I) Alcohol and Quality

In [20]:

```
1 fig = plt.figure(figsize = (10,6))
2 sns.barplot(x = 'quality', y = 'alcohol', data = wine_data)
3 '''Alcohol content increases as the quality of wine increases'''
```

Out[20]:

'Alcohol content increases as the quality of wine increases'



From the above analysis it is assumed that the following variables have higher importance in determining the quality of the wine

- volatile acidity
- citric acid
- chlorides
- sulphates
- pH
- alcohol

## Checking for Correlation among the variables




In [21]:

```

1 # Plotting the relationship between variables
2 plt.figure(figsize=(8,8))
3 corr_matrix = wine_data.corr()
4 cols = corr_matrix.nlargest(12,'quality')['quality'].index
5 cm = np.corrcoef(wine_data[cols].values.T)
6 sns.set(font_scale=1.25)
7 heatmap = sns.heatmap(corr_matrix,annot=True,fmt='.2f',annot_kws={'size': 10},cbar=True)
8 plt.show()

```



There is a fairly positive correlation between:

- quality and pH (0.67)
- quality and sulphates (0.67)
- residual sugar and free sulfur dioxide (0.67)

There is a fairly negative correlation between:

- quality and chlorides (-0.68)

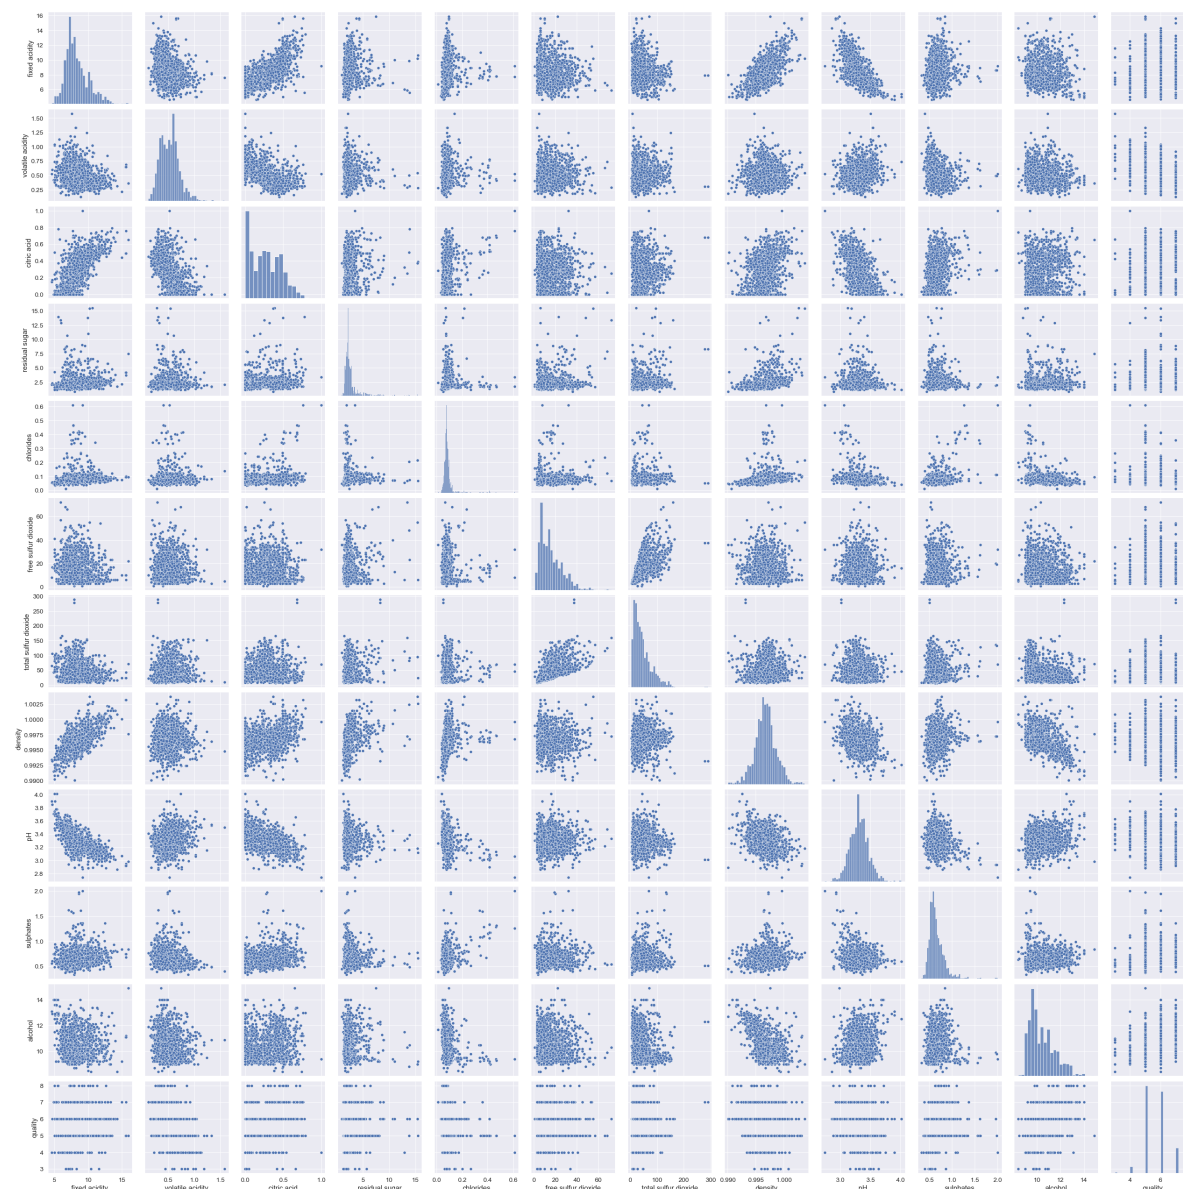
- alcohol and sulphates (-0.55)
- sulphates and chlorides (-0.54)
- pH and total sulfur dioxide (-0.50)

In [22]:

```
1 # Pair plots
2 sns.pairplot(wine_data,height=3.0)
```

Out[22]:

&lt;seaborn.axisgrid.PairGrid at 0x21ae2f2cc40&gt;



## Outlier Detection

### Plotting the important variables

#### a) volatile acidity


In [23]:



```
1 sns.boxplot(x = 'volatile acidity', data = wine_data)
```

Out[23]:

&lt;AxesSubplot:xlabel='volatile acidity'&gt;



In [24]:



```
1 '''There are outliers present in volatile acidity'''
```

Out[24]:

```
'There are outliers present in volatile acidity'
```


**b) citric acid**

In [25]:

```
1 sns.boxplot(x = 'citric acid', data = wine_data)
2 '''There are very few outliers in the data'''
```

Out[25]:

'There are very few outliers in the data'



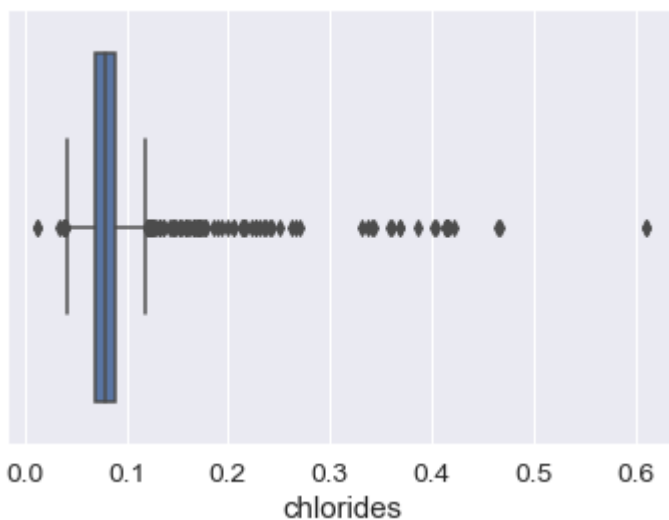
### c) chlorides

In [26]:

```
1 sns.boxplot(x = 'chlorides', data = wine_data)
2 '''There are many outliers in the data'''
```

Out[26]:

'There are many outliers in the data'




### d) sulphates

In [27]:

```
1 sns.boxplot(x = 'sulphates', data = wine_data)
2 '''There are many outliers in the data'''
```

Out[27]:

'There are many outliers in the data'




e) pH

In [28]:

```
1 sns.boxplot(x = 'pH', data = wine_data)
2 '''There are some outliers in the data'''
```

Out[28]:

'There are some outliers in the data'




f) alcohol

In [29]:

```
1 sns.boxplot(x = 'alcohol', data = wine_data)
2 '''There are very few outliers in the data'''
```

Out[29]:

'There are very few outliers in the data'



## Data Preprocessing

In [30]:

```
1 # Taking a copy of the wine_data. Preprocessing is done on the copy
2 wine_data_preprocessed = wine_data.copy()
3 wine_data_preprocessed.tail()
```

Out[30]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	al
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	

## Removing highly correlated variables

There is a fairly strong correlation between "free sulfur dioxide" and "residual sugar". During analysis, it has been found that "residual sugar" does not vary significantly based on quality of the wine. Hence we can remove this column from our dataset.

In [31]:



```

1 # Removing the highly correlated variable
2 wine_data_preprocessed.drop(columns='residual sugar',inplace=True)
3 wine_data_preprocessed.head()

```

Out[31]:

	fixed acidity	volatile acidity	citric acid	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5

## Adding new column 'wine rating' as the target variable

### Assumption

In the given data set, the column 'quality' has values [3,4,5,6,7,8] indicating the quality of the wine. To enable us to create models that can rate the wine as Bad or Good, we make the following assumptions:

- if the value of 'quality' is 7 or 8 the wine is Good
- if the 'quality' of wine is 3,4,5 or 6, the wine is Bad
- 'Good' is represented by 1 while 'Bad' by 0

In [32]:



```

1 wine_data_preprocessed['wine rating']=[1 if x>=7 else 0 for x in wine_data_preprocessed['quality']]
2 wine_data_preprocessed[wine_data_preprocessed['wine rating'] == 1].tail()

```

Out[32]:

	fixed acidity	volatile acidity	citric acid	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	qu
1541	7.4	0.25	0.29	0.054	19.0	49.0	0.99666	3.40	0.76	10.90	
1544	8.4	0.37	0.43	0.063	12.0	19.0	0.99550	3.17	0.81	11.20	
1549	7.4	0.36	0.30	0.074	17.0	24.0	0.99419	3.24	0.70	11.40	
1555	7.0	0.56	0.17	0.065	15.0	24.0	0.99514	3.44	0.68	10.55	
1584	6.7	0.32	0.44	0.061	24.0	34.0	0.99484	3.29	0.80	11.60	

In [33]:

```
1 wine_data_preprocessed[wine_data_preprocessed['wine rating'] == 0]
```

Out[33]:

	fixed acidity	volatile acidity	citric acid	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	qu
0	7.4	0.700	0.00	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
1	7.8	0.880	0.00	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	
2	7.8	0.760	0.04	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	
3	11.2	0.280	0.56	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	
4	7.4	0.700	0.00	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	
...	...	...	...	...	...	...	...	...	...	...	...
1594	6.2	0.600	0.08	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	
1595	5.9	0.550	0.10	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	
1596	6.3	0.510	0.13	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	
1597	5.9	0.645	0.12	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	
1598	6.0	0.310	0.47	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	

1382 rows × 12 columns

## Outlier Removal

In [34]:

```
1 # z-score method of Outlier Removal
2 z_scores = zscore(wine_data_preprocessed)
3
4 # taking the absolute values of the z-score
5 abs_z_scores = np.abs(z_scores)
6
7 # filter condition for z_score - An outlier of a dataset is defined as a value that is
8 filtered_entries = (abs_z_scores < 3).all(axis=1)
9
10 # populating the data frame with data having only absolute z_score < 3
11 wine_data_preprocessed = wine_data_preprocessed[filtered_entries]
12
13 print(f"After outlier removal {wine_data_preprocessed.shape[0]} rows and {wine_data_preprocessed.shape[1]} columns are left")
14
```

After outlier removal 1469 rows and 12 columns are left

## Detection and Removal of Multi-Collinearity using Variance Inflation Factor(VIF)



In [35]:



```
1 wine_data_preprocessed.columns
```

Out[35]:

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'chlorides',  
      'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH',  
      'sulphates', 'alcohol', 'quality', 'wine rating'],  
      dtype='object')
```

In [36]:



```

1  # Multicollinearity occurs when there are two or more independent variables in a regres
2  # When the features are highly collinear, it becomes difficult to determine their indiv
3  # Here we are using VIF method to determine the multi-collinearity
4
5  # Idea of VIF method
6  # In VIF method, it takes one column at a time as target and others as features and fit
7  # After this, it calculates the Rsquare value and for the VIF value, we take the invers
8  # Hence after each iteration, we get VIF value for each column (which was taken as targ
9
10 # Higher the VIF, greater is the correlation between the variables
11 # VIF exceeding 5 or 10 indicates high multicollinearity between this independent vari
12
13 #Assumption
14 # Here we are considering only cases where VIF is greater than 10. If VIF is greater th
15
16 # All the independent variable
17 variables = wine_data_preprocessed[['fixed acidity', 'volatile acidity', 'citric acid',
18                                     'chlorides', 'free sulfur dioxide',
19                                     'total sulfur dioxide', 'density', 'pH', 'sulphates', 'alcohol',]]
20
21 # adding an intercept
22 x = add_constant(variables)
23 vif = pd.DataFrame()
24
25 # calculating the VIF for each variable
26 vif['VIF'] = [variance_inflation_factor(x.values,i) for i in range (x.shape[1])]
27 vif['features'] = x.columns
28 vif.round(1)
29

```

Out[36]:

	VIF	features
0	1356727.9	const
1	6.5	fixed acidity
2	1.9	volatile acidity
3	3.1	citric acid
4	1.2	chlorides
5	1.9	free sulfur dioxide
6	2.2	total sulfur dioxide
7	4.2	density
8	2.9	pH
9	1.3	sulphates
10	2.3	alcohol

None of the columns have VIF > 10, so we are not dropping any columns

## Splitting the data into dependent and independent variables

In [37]:

```
1 # independent variables
2 X = wine_data_preprocessed.drop(['quality', 'wine rating'], axis=1)
3 # dependent variable
4 Y = wine_data_preprocessed['wine rating']
5 print('The independent variables: ')
6 print(X.columns)
7 print("")
8 print("")
9 print('The dependent variable: ')
10 print(Y.name)
11
```

The independent variables:

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'chlorides',
      'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH',
      'sulphates', 'alcohol'],
      dtype='object')
```

The dependent variable:

wine rating

## Splitting the data into Training set and Testing set

In [38]:

```
1 """Train-Test Split is 80:20 i.e. 80% training data and 20% testing data"""
```

Out[38]:

```
'Train-Test Split is 80:20 i.e. 80% training data and 20% testing data'
```

In [39]:

```
1 #Train and Test splitting of data
2 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state=42)
```

## Standardizing the data

In [40]:

```
1 # Scaling the train and test data
2 scaler = StandardScaler()
3 X_train = scaler.fit_transform(X_train)
4 X_test = scaler.transform(X_test)
```

## Model Building

In [41]:

```

1  ## For comparing Different Models
2  model_comparison_df = pd.DataFrame(columns = ['model', 'accuracy', 'error_rate', 'auc'])
3  model_comparison_df.set_index('model', inplace=True)
4  model_comparison_df

```

Out[41]:

	accuracy	error_rate	auc
model			

## Decision Tree Classifier

In [42]:

```

1  # Creating a DecisionTreeClassifier with default values of hyper parameters
2  model_dt = DecisionTreeClassifier()
3  print("Deafult Parameters:\n\n\n" + str(model_dt.get_params(True)))

```

Deafult Parameters:

```

{'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_depth': N
one, 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease':
0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'min_samples_split':
2, 'min_weight_fraction_leaf': 0.0, 'presort': 'deprecated', 'random_state':
None, 'splitter': 'best'}

```

## Training the Model and Prediction

In [43]:

```

1
2  # trains the Decision Tree Model using the training sets(x_train,y_train)
3  model_dt.fit(X_train,Y_train)
4
5  # calculation of score(accuracy) of training
6  training_score = model_dt.score(X_train,Y_train)
7  print("Training Score: " + str(training_score))
8
9  # checking the score of testing
10 testing_score = model_dt.score(X_test,Y_test)
11 print("Testing Score: " + str(testing_score))

```

Training Score: 1.0

Testing Score: 0.8741496598639455

**This is a case of overfitting as the Training Score is much higher than Testing score**

## Hyper Parameter Tuning using Grid Search CV

In [44]:



```

1 #predefined set of hyperparameters
2 paramlist = {'max_depth':range(1,15),
3             'min_samples_split':range(2,10),
4             'min_samples_leaf':range(20,51,10),
5             'criterion': ['gini','entropy']}
6
7 # creating a DecisionTreeClassifierModel
8 decision_model = DecisionTreeClassifier(random_state=0)
9
10
11 start = time.time()
12 # creating an object of GridSearchCV
13 grid_cv = GridSearchCV(decision_model,cv=10,param_grid=paramlist)
14
15 # training the model
16 grid_cv.fit(X_train,Y_train)
17 end = time.time()
18
19 print('Best Parameters for Decision Tree Classifier using Grid Search CV : \n', grid_cv
20 print('Time taken in grid search: {0: .2f}'.format(end - start))

```

Best Parameters for Decision Tree Classifier using Grid Search CV :  
 {'criterion': 'gini', 'max\_depth': 3, 'min\_samples\_leaf': 50, 'min\_samples\_split': 2}  
 Time taken in grid search: 64.28

### Creating Decision Tree based on result of GridSearchCV - CART

In [45]:



```

1 # Instatiating the CART model
2 cart_md1 = DecisionTreeClassifier(random_state=0,max_depth=3,min_samples_leaf=50,min_s
3 print("Parameters for CART model :\n\n\n" + str(cart_md1.get_params(True)))
4
5 # Adding a row for Decision Tree Classifier in the model_comparison table
6 model_comparison_df = model_comparison_df.append({'model': 'Decision Tree Classifier'},

```

Parameters for CART model :

```

{'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_depth':
3, 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease': 0.
0, 'min_impurity_split': None, 'min_samples_leaf': 50, 'min_samples_split':
2, 'min_weight_fraction_leaf': 0.0, 'presort': 'deprecated', 'random_state':
0, 'splitter': 'best'}

```

### Training the Model and Prediction

In [46]:



```

1 # fitting the model with Training set
2 cart_md1.fit(X_train,Y_train)
3
4 # calculation of Training score
5 training_score = cart_md1.score(X_train,Y_train)
6 print("Training score :" + str(training_score))
7
8 # calculation of Test score
9 testing_score = cart_md1.score(X_test,Y_test)
10 print("Testing score :" + str(testing_score))

```

Training score :0.8910638297872341

Testing score :0.8877551020408163

**The Training Score and Testing score are statistically not very different**

## Evaluating the model

### a) Classification Report

In [47]:



```

1 # predicting the values for X-test
2 prediction_cart= cart_md1.predict(X_test)
3
4 cart_report = classification_report(Y_test, prediction_cart,output_dict=True)
5
6 print(classification_report(Y_test, prediction_cart))

```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	252
1	0.76	0.31	0.44	42
accuracy			0.89	294
macro avg	0.83	0.65	0.69	294
weighted avg	0.88	0.89	0.87	294

**Accuracy of the Decision Tree Classifier (CART) is 89%**

### b) Confusion Matrix

In [48]:

```

1 def create_conf_mat(test_class_set, predictions):
2     """Function returns confusion matrix comparing two arrays"""
3     if (len(test_class_set.shape) != len(predictions.shape) == 1):
4         return print('Arrays entered are not 1-D.\nPlease enter the correctly sized set')
5     elif (test_class_set.shape != predictions.shape):
6         return print('Number of values inside the Arrays are not equal to each other.\n')
7     else:
8         # Set Metrics
9         test_crosstb_comp = pd.crosstab(index = test_class_set,
10                                         columns = predictions)
11         # Changed for Future deprecation of as_matrix
12         test_crosstb = test_crosstb_comp.values
13         return test_crosstb


```

In [49]:

```

1 conf_mat = create_conf_mat(Y_test, prediction_cart)
2 sns.heatmap(conf_mat, annot=True, fmt='d', cbar=False)
3 plt.xlabel('Predicted Values')
4 plt.ylabel('Actual Values')
5 plt.title('Actual vs. Predicted Confusion Matrix')
6 plt.show()

```



### c) Accuracy

In [50]:



```

1 accuracy_cart = cart_md1.score(X_test, Y_test)
2
3 print("Mean accuracy on the test set (when using Decision Tree Classifier (CART)):\n {0: .4f}"
4       .format(accuracy_cart))
5
6 ## adding the accuracy to the comparison table
7 model_comparison_df.loc[[0], 'accuracy'] = accuracy_cart
8 model_comparison_df

```

Mean accuracy on the test set (when using Decision Tree Classifier (CART)):  
0.888

Out[50]:

	accuracy	error_rate	auc	model
0	0.887755	NaN	NaN	Decision Tree Classifier

#### d) Error Rate

In [51]:



```

1 test_error_rate_cart = 1 - accuracy_cart
2 print("The test error rate for the Decision Tree (CART) Classifier model is:\n {0: .4f}"
3       .format(test_error_rate_cart))
4
5 ## adding the error rate to the comparison table
6 model_comparison_df.loc[[0], 'error_rate'] = test_error_rate_cart
7 model_comparison_df

```

The test error rate for the Decision Tree (CART) Classifier model is:  
0.1122

Out[51]:

	accuracy	error_rate	auc	model
0	0.887755	0.112245	NaN	Decision Tree Classifier

#### e) Area under the curve





In [52]:

```

1 def plot_roc_curve(fpr, tpr, auc, estimator, xlim=None, ylim=None):
2     """
3     Purpose
4     -----
5     Function creates ROC Curve for respective model given selected parameters.
6     Optional x and y limits to zoom into graph
7
8     Parameters
9     -----
10    * fpr: Array returned from sklearn.metrics.roc_curve for increasing
11          false positive rates
12    * tpr: Array returned from sklearn.metrics.roc_curve for increasing
13          true positive rates
14    * auc: Float returned from sklearn.metrics.auc (Area under Curve)
15    * estimator: String representation of appropriate model, can only contain the
16    following: ['knn', 'rf', 'dt', 'lr']
17    * xlim: Set upper and lower x-limits
18    * ylim: Set upper and lower y-limits
19    """
20    my_estimators = {'knn': ['Kth Nearest Neighbor', 'deeppink'],
21                    'rf': ['Random Forest', 'red'],
22                    'dt': ['Decision Tree', 'blue'],
23                    'lr': ['Logistic Regression', 'purple']}
24
25    try:
26        plot_title = my_estimators[estimator][0]
27        color_value = my_estimators[estimator][1]
28    except KeyError as e:
29        print("{}' does not correspond with the appropriate key inside the estimators
30 \nPlease refer to function to check `my_estimators` dictionary.".format(estimator))
31        raise
32
33    fig, ax = plt.subplots(figsize=(10, 10))
34    ax.set_facecolor('#fafafa')
35
36    plt.plot(fpr, tpr,
37             color=color_value,
38             linewidth=1)
39    plt.title('ROC Curve For {0} (AUC = {1: 0.3f})'\
40             .format(plot_title, auc))
41
42    plt.plot([0, 1], [0, 1], 'k--', lw=2) # Add Diagonal line
43    plt.plot([0, 0], [1, 0], 'k--', lw=2, color = 'black')
44    plt.plot([1, 0], [1, 1], 'k--', lw=2, color = 'black')
45    if xlim is not None:
46        plt.xlim(*xlim)
47    if ylim is not None:
48        plt.ylim(*ylim)
49    plt.xlabel('False Positive Rate')
50    plt.ylabel('True Positive Rate')
51    plt.show()
52    plt.close()

```

In [53]:



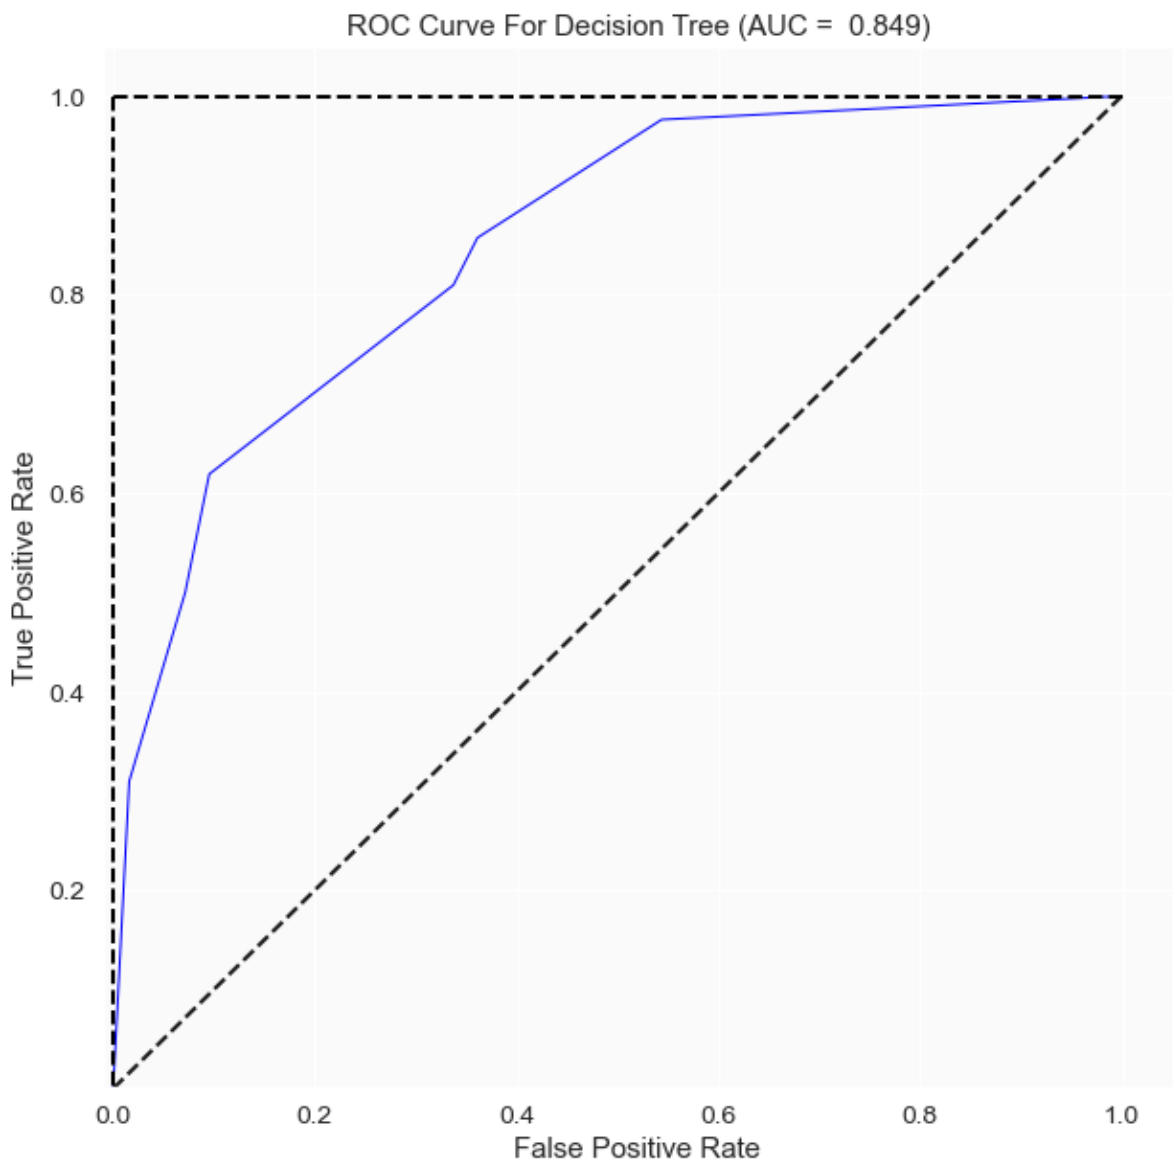
```
1 predictions_prob_cart = cart_mdl.predict_proba(X_test)[: , 1]
2
3 fpr_cart, tpr_cart, _ = roc_curve(Y_test,
4                                   predictions_prob_cart,
5                                   pos_label = 1)
6
7 # calculating area under the tpr-fpr curve
8 auc_cart = auc(fpr_cart, tpr_cart)
9
10 ## adding the auc to the comparison table
11 model_comparison_df.loc[[0], 'auc'] = auc_cart
12 model_comparison_df
```

Out[53]:

	accuracy	error_rate	auc	model
0	0.887755	0.112245	0.849254	Decision Tree Classifier

In [54]:

```
1 # plotting the roc curve
2 plot_roc_curve(fpr_cart, tpr_cart, auc_cart, 'dt',
3               xlim=(-0.01, 1.05),
4               ylim=(0.001, 1.05))
```



The Area under the curve for the TPR-FPR curve is about 84.9%

In [55]:

```
1 ## adding the precision and recall for Decision Tree Classifier (CART) to the comparison
2 model_comparison_df.loc[[0], 'precision'] = metrics.precision_score(Y_test, prediction_cart)
3 model_comparison_df.loc[[0], 'recall'] = metrics.recall_score(Y_test, prediction_cart)
4 model_comparison_df
```

Out[55]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524

## Random Forest Classifier

In [56]:

```
1 # Creating a Random Forest Classifier with default values of hyper parameters
2 rf_model = RandomForestClassifier(random_state=0)
3 print("Deafult Parameters:\n\n" + str(rf_model.get_params(True)))
```

Deafult Parameters:

```
{'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_depth': None, 'max_features': 'auto', 'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease': 0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_estimators': 100, 'n_jobs': None, 'oob_score': False, 'random_state': 0, 'verbose': 0, 'warm_start': False}
```

## Training the Model and Prediction

In [57]:

```
1 # training the Random Forest Classifier
2 rf_model.fit(X_train, Y_train)
3
4
5 # calculation of Training score
6 training_score_rf = rf_model.score(X_train,Y_train)
7 print("Training score :" + str(training_score_rf))
8
9 # calculation of Test score
10 testing_score_rf = rf_model.score(X_test,Y_test)
11 print("Testing score :" + str(testing_score_rf))
12
13 # predicting the values for X-test
14 prediction_rf= rf_model.predict(X_test)
15
```

Training score :1.0

Testing score :0.9081632653061225

Since the training score is higher than testing score, it indicates an overfit

## Hyperparameter Optimization of RandomForest using GridSearchCV

In [58]:



```

1 # creating an instance of RandomForestClassifier
2 fit_rf = RandomForestClassifier(random_state=40)
3 np.random.seed(0)
4 start = time.time()
5
6 # predefined set of parameters for GridSearchCV
7 param_dist = {'max_depth':[2,3,4,5],
8               'bootstrap':[True,False],
9               'max_features':['auto','sqrt','log2',None],
10              'criterion':['gini','entropy']}
11
12 cv_rf = GridSearchCV(fit_rf,cv=10,param_grid=param_dist,n_jobs=3)
13
14 cv_rf.fit(X_train,Y_train)
15 print('Best Parameters for Random Forest using Grid search: \n',cv_rf.best_params_)
16 end = time.time()
17 print('Time taken in grid search: %0.2f'%(end-start))

```

Best Parameters for Random Forest using Grid search:

```
{'bootstrap': True, 'criterion': 'gini', 'max_depth': 5, 'max_features': None}
```

Time taken in grid search: 98.38

In [59]:



```

1 # setting the model with the best paramters suggested by GridSearchCV
2 fit_rf.set_params(bootstrap=True,criterion='gini',max_features = None,max_depth = 5)

```

Out[59]:

RandomForestClassifier(max\_depth=5, max\_features=None, random\_state=40)

### Determining the number of trees using OOB rate

In [60]:



```

1 fit_rf.set_params(warm_start=True,
2                   oob_score=True)
3
4 min_estimators = 100
5 max_estimators = 1000
6
7 error_rate = {}
8
9 for i in range(min_estimators, max_estimators + 1):
10     fit_rf.set_params(n_estimators=i)
11     fit_rf.fit(X_train, Y_train)
12
13     oob_error = 1 - fit_rf.oob_score_
14     error_rate[i] = oob_error

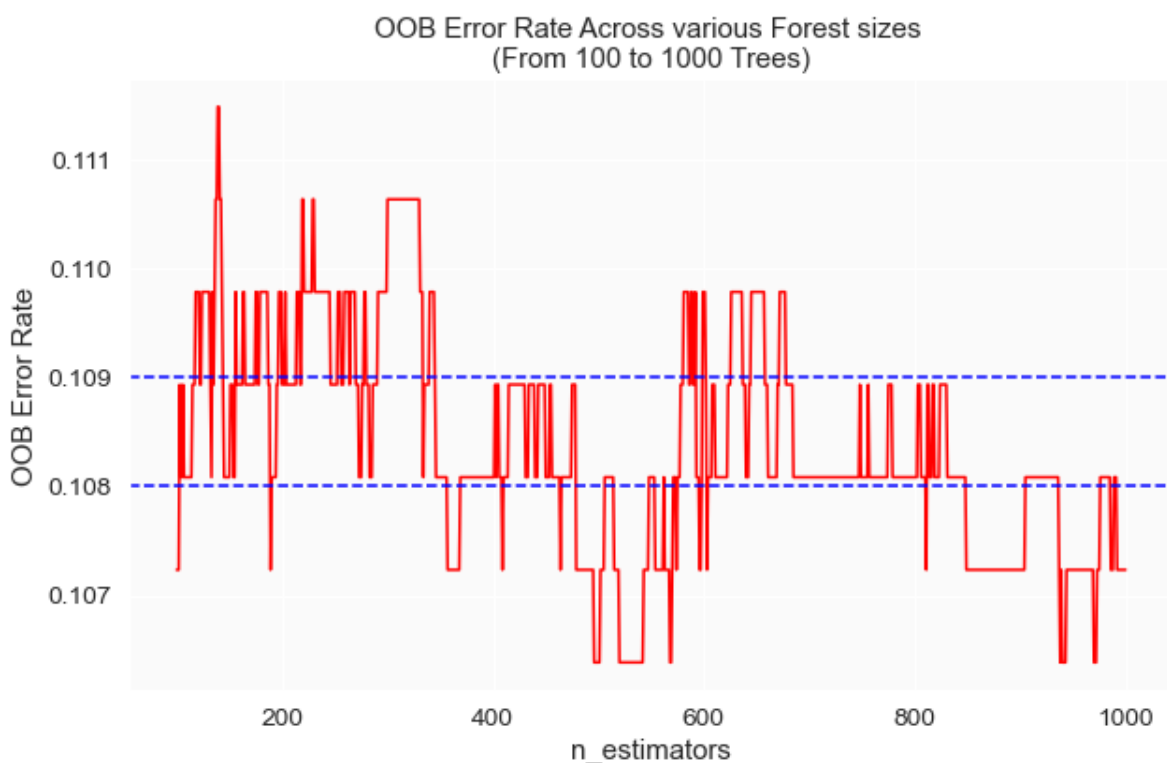
```

In [61]:

```
1 # Plotting the oob rate
2 oob_series = pd.Series(error_rate)
3 fig, ax = plt.subplots(figsize=(10,6))
4
5 ax.set_facecolor('#fafafa')
6
7 oob_series.plot(kind='line',color = 'red')
8 plt.axhline(0.109,color='blue',linestyle='--')
9 plt.axhline(0.108,color='blue',linestyle='--')
10 plt.xlabel('n_estimators')
11 plt.ylabel('OOB Error Rate')
12 plt.title('OOB Error Rate Across various Forest sizes \n(From 100 to 1000 Trees)')
```

Out[61]:

Text(0.5, 1.0, 'OOB Error Rate Across various Forest sizes \n(From 100 to 1000 Trees)')



In [62]:

```
1 print('OOB Error rate for 400 trees is: {0:.5f}'.format(oob_series[400]))
```

OOB Error rate for 400 trees is: 0.10809

**When  $n_{\text{estimators}} = 400$ , the error rate relatively remains stable so we choose  $n_{\text{estimators}} = 400$**

In [63]:

```

1 # fitting the random forest with n_estimators
2 fit_rf.set_params(n_estimators=400,bootstrap=True,warm_start=False,oob_score=False)
3 # fit_rf.set_params(bootstrap=True,criterion='gini',max_features = None,max_depth = 5)
4 print("The parameters of the random forest are: \n\n" + str(fit_rf.get_params(True)))

```

The parameters of the random forest are:

```
{'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_depth': 5, 'max_features': None, 'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease': 0.0, 'min_impurity_split': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_estimators': 400, 'n_jobs': None, 'oob_score': False, 'random_state': 40, 'verbose': 0, 'warm_start': False}
```

## Training the model

In [64]:

```

1 fit_rf.fit(X_train,Y_train)
2
3 # calculation of Training score
4 training_score_rf = fit_rf.score(X_train,Y_train)
5 print("Training score :" + str(training_score_rf))
6
7 # calculation of Test score
8 testing_score_rf = fit_rf.score(X_test,Y_test)
9 print("Testing score :" + str(testing_score_rf))
10
11 # predicting the values for X-test
12 prediction_rf= fit_rf.predict(X_test)
13
14 # Adding a row for Random Forest Classifier in the model_comparison table
15 model_comparison_df = model_comparison_df.append({'model': 'Random Forest Classifier'},
16 model_comparison_df

```

Training score :0.9446808510638298

Testing score :0.9047619047619048

Out[64]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	NaN	NaN	NaN	Random Forest Classifier	NaN	NaN

## Evaluating the model

### a) Classification Report

In [65]:



```
1 print(classification_report(Y_test, prediction_rf))
```

	precision	recall	f1-score	support
0	0.91	0.98	0.95	252
1	0.79	0.45	0.58	42
accuracy			0.90	294
macro avg	0.85	0.72	0.76	294
weighted avg	0.90	0.90	0.89	294


**Accuracy of the Random Forest Classifier is 90%**

## b) Confusion Matrix

In [66]:



```
1 conf_mat = create_conf_mat(Y_test, prediction_rf)
2 sns.heatmap(conf_mat, annot=True, fmt='d', cbar=False)
3 plt.xlabel('Predicted Values')
4 plt.ylabel('Actual Values')
5 plt.title('Actual vs. Predicted Confusion Matrix')
6 plt.show()
```



## c) Accuracy



In [67]:



```

1 accuracy_rf = fit_rf.score(X_test, Y_test)
2
3 print("Mean accuracy on the test set (when using Random Forest Classifier):\n {0:.3f}"
4       .format(accuracy_rf))
5
6 ## adding the accuracy to the comparison table
7 model_comparison_df.loc[[1], 'accuracy'] = accuracy_rf
8 model_comparison_df

```

Mean accuracy on the test set (when using Random Forest Classifier):  
0.905

Out[67]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	NaN	NaN	Random Forest Classifier	NaN	NaN

#### d) Error Rate

In [68]:



```

1 test_error_rate_rf = 1 - accuracy_rf
2 print("The test error rate for the Random Forest Classifier model is:\n {0: .4f}"
3       .format(test_error_rate_rf))
4
5 ## adding the error_rate to the comparison table
6 model_comparison_df.loc[[1], 'error_rate'] = test_error_rate_rf
7 model_comparison_df

```

The test error rate for the Random Forest Classifier model is:  
0.0952

Out[68]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	NaN	Random Forest Classifier	NaN	NaN

#### e) Area Under the Curve

In [69]:




```
1 predictions_prob_rf = fit_rf.predict_proba(X_test)[:, 1]
2
3 fpr_rf, tpr_rf, _ = roc_curve(Y_test,
4                               predictions_prob_rf,
5                               pos_label = 1)
6
7 auc_rf = auc(fpr_rf, tpr_rf)
8 ## adding the auc to the comparison table
9 model_comparison_df.loc[[1], 'auc'] = auc_rf
10 model_comparison_df
```

Out[69]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	NaN	NaN

In [70]:

```
1 plot_roc_curve(fpr_rf, tpr_rf, auc_rf, 'rf',
2               xlim=(-0.01, 1.05),
3               ylim=(0.001, 1.05))
```



The Area under the curve for the TPR-FPR curve is about 91.3%

In [71]:

```
1 ## adding the precision and recall for Random Forest to the comparison table
2 model_comparison_df.loc[[1], 'precision'] = metrics.precision_score(Y_test, prediction_rf)
3 model_comparison_df.loc[[1], 'recall'] = metrics.recall_score(Y_test, prediction_rf)
4 model_comparison_df
```

Out[71]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381

## Visualizing the Feature Importance

In [72]:



```
1 def variable_importance(fit):
2     """
3     Purpose
4     -----
5     Checks if model is fitted CART model then produces variable importance
6     and respective indices in dictionary.
7
8     Parameters
9     -----
10    * fit: Fitted model containing the attribute feature_importances_
11
12    Returns
13    -----
14    Dictionary containing arrays with importance score and index of columns
15    ordered in descending order of importance.
16    """
17    try:
18        if not hasattr(fit, 'fit'):
19            return print("'{}' [is not an instantiated model from scikit-learn]".format(fit))
20
21        # Captures whether the model has been trained
22        if not vars(fit)["estimators_"]:
23            return print("Model does not appear to be trained.")
24    except KeyError:
25        print("Model entered does not contain 'estimators_' attribute.")
26
27    importances = fit.feature_importances_
28    indices = np.argsort(importances)[::-1]
29    return {'importance': importances,
30            'index': indices}
```

In [73]:



```

1 def variable_importance_plot(importance, indices, name_index):
2     """
3     Purpose
4     -----
5     Prints bar chart detailing variable importance for CART model
6     NOTE: feature_space list was created because the bar chart
7     was transposed and index would be in incorrect order.
8
9     Parameters
10    -----
11    * importance: Array returned from feature_importances_ for CART
12                  models organized by dataframe index
13    * indices: Organized index of dataframe from largest to smallest
14               based on feature_importances_
15    * name_index: Name of columns included in model
16
17    Returns:
18    -----
19    Returns variable importance plot in descending order
20    """
21    index = np.arange(len(name_index))
22
23    importance_desc = sorted(importance)
24    feature_space = []
25    for i in range(indices.shape[0] - 1, -1, -1):
26        feature_space.append(name_index[indices[i]])
27
28    fig, ax = plt.subplots(figsize=(6,6))
29
30    # ax.set_axis_bgcolor('#fafafa')
31    plt.title('Feature importances for Random Forest Model\
32    \nWine Quality Prediction')
33    plt.barh(index,
34             importance_desc,
35             align="center",
36             color = '#875FDB')
37    plt.yticks(index,
38              feature_space)
39
40    plt.ylim(-1, 10)
41    plt.xlim(0, max(importance_desc) + 0.01)
42    plt.xlabel('Mean Decrease in Impurity')
43    plt.ylabel('Feature')
44
45    plt.show()
46    plt.close()

```

In [74]:



```

1 var_imp_rf = variable_importance(fit_rf)
2
3 importances_rf = var_imp_rf['importance']
4
5 indices_rf = var_imp_rf['index']
6 print(indices_rf.shape)

```

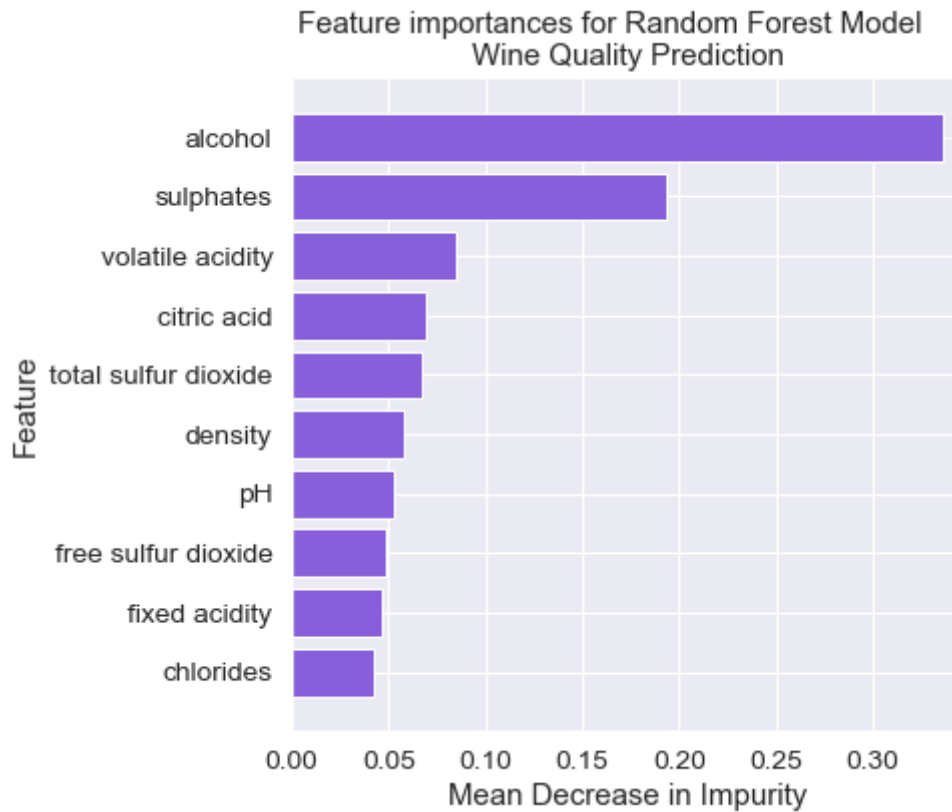
(10,)

In [75]:

```
1 names_index = ['fixed acidity', 'volatile acidity', 'citric acid', 'chlorides',  
2               'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH',  
3               'sulphates', 'alcohol']
```

In [76]:

```
1 variable_importance_plot(importances_rf, indices_rf, names_index)
```



The three main features that determine the quality of wine are : alcohol, sulphates and density

## Logistic Regression

In [77]:



```

1 # Creating an instance of Logistic Regression with default values of hyper parameters
2 lr_model = LogisticRegression(random_state=0)
3 print("Deafult Parameters:\n\n\n" + str(lr_model.get_params(True)))
4
5 # Adding a row for Logistic Regression in the model_comparison table
6 model_comparison_df = model_comparison_df.append({'model': 'Logistic Regression'}, ignore_index=True)
7 model_comparison_df

```

Deafult Parameters:

```

{'C': 1.0, 'class_weight': None, 'dual': False, 'fit_intercept': True, 'intercept_scaling': 1, 'l1_ratio': None, 'max_iter': 100, 'multi_class': 'auto', 'n_jobs': None, 'penalty': 'l2', 'random_state': 0, 'solver': 'lbfgs', 'tol': 0.0001, 'verbose': 0, 'warm_start': False}

```

Out[77]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381
2	NaN	NaN	NaN	Logistic Regression	NaN	NaN

## Training the model and Prediction

In [78]:



```

1 # fitting the Logistic Regression model with the training data and training output
2 lr_model.fit(X_train, Y_train)
3
4 # predicting the values for X-test
5 prediction_lr= lr_model.predict(X_test)

```

## Evaluating the model

### a) Classification Report

In [79]:



```
1 print(classification_report(Y_test, prediction_lr))
2
```

	precision	recall	f1-score	support
0	0.91	0.97	0.94	252
1	0.71	0.40	0.52	42
accuracy			0.89	294
macro avg	0.81	0.69	0.73	294
weighted avg	0.88	0.89	0.88	294


**Accuracy of the Logistic Regression Model is 89%**

## b) Confusion Matrix

In [80]:



```
1 conf_mat = create_conf_mat(Y_test, prediction_lr)
2 sns.heatmap(conf_mat, annot=True, fmt='d', cbar=False)
3 plt.xlabel('Predicted Values')
4 plt.ylabel('Actual Values')
5 plt.title('Actual vs. Predicted Confusion Matrix')
6 plt.show()
```



## c) Accuracy



In [81]:



```

1 accuracy_lr = lr_model.score(X_test, Y_test)
2
3 print("Mean accuracy on the test set (when using Logistic Regression):\n {0:.3f}" \
4       .format(accuracy_lr))
5
6 ## adding the auc to the comparison table
7 model_comparison_df.loc[[2], 'accuracy'] = accuracy_lr
8 model_comparison_df

```

Mean accuracy on the test set (when using Logistic Regression):  
0.891

Out[81]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381
2	0.891156	NaN	NaN	Logistic Regression	NaN	NaN

#### d) Error Rate

In [82]:



```

1 test_error_rate_lr = 1 - accuracy_lr
2 print("The test error rate for the Logistic Regression model is:\n {0: .4f}" \
3       .format(test_error_rate_lr))
4
5 ## adding the error_rate to the comparison table
6 model_comparison_df.loc[[2], 'error_rate'] = test_error_rate_lr
7 model_comparison_df

```

The test error rate for the Logistic Regression model is:  
0.1088

Out[82]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381
2	0.891156	0.108844	NaN	Logistic Regression	NaN	NaN

#### e) Area under the curve

In [83]:



```

1 predictions_prob_lr = lr_model.predict_proba(X_test)[:, 1]
2
3 # calculating the FPR and TPR
4 fpr_lr, tpr_lr, _lr = roc_curve(Y_test,
5                               predictions_prob_lr,
6                               pos_label = 1)
7
8 # calculating the area under the curve
9 auc_lr = auc(fpr_lr, tpr_lr)
10
11 ## adding the auc to the comparison table
12 model_comparison_df.loc[[2], 'auc'] = auc_lr
13 model_comparison_df

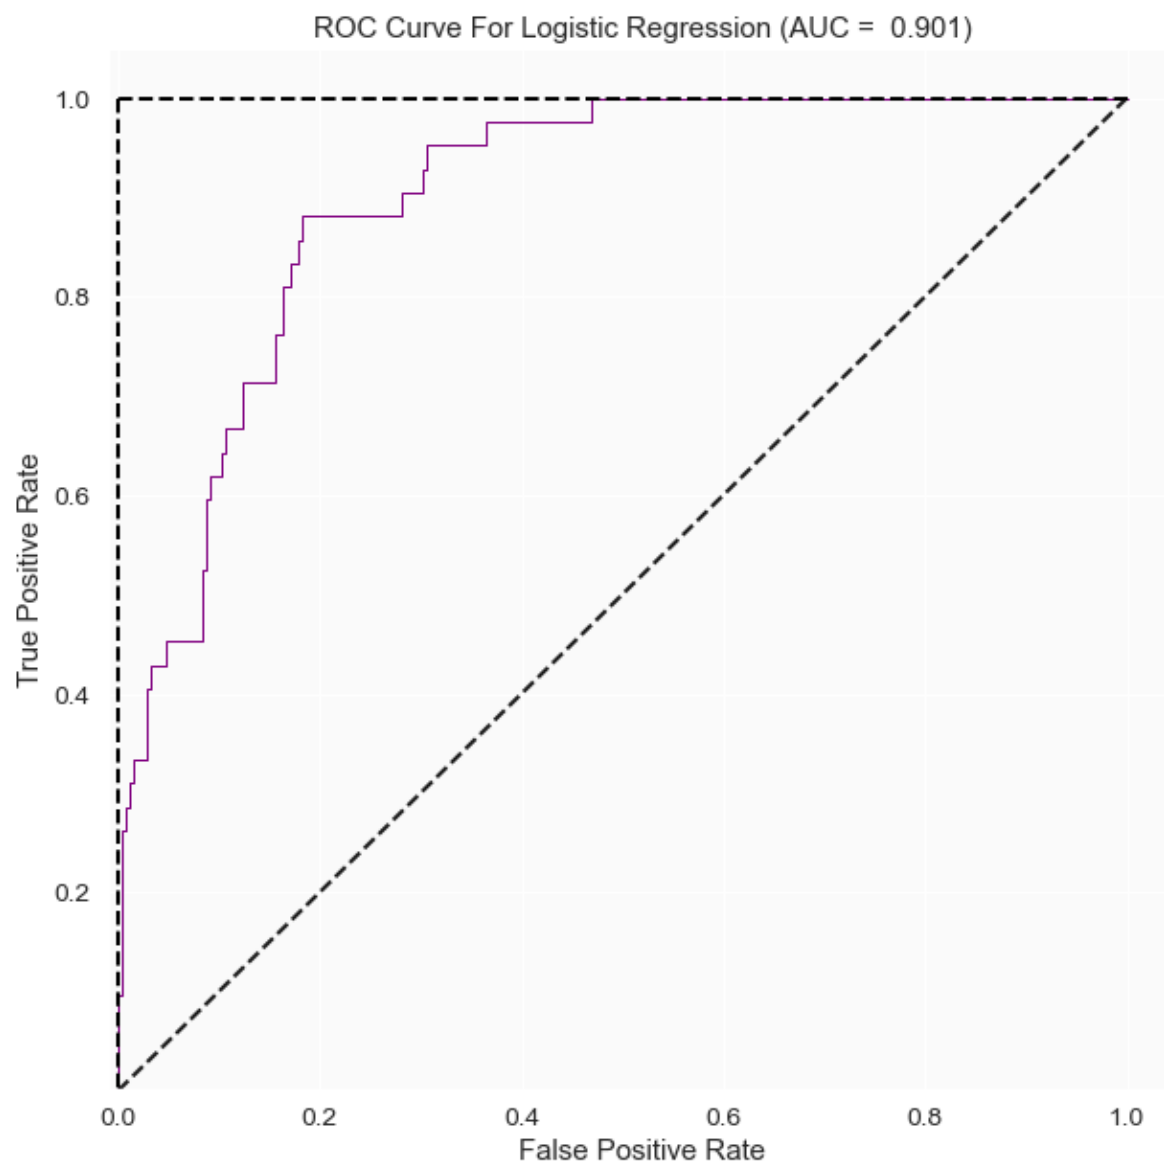
```

Out[83]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381
2	0.891156	0.108844	0.900605	Logistic Regression	NaN	NaN

In [84]:

```
1 plot_roc_curve(fpr_lr, tpr_lr, auc_lr, 'lr',  
2                 xlim=(-0.01, 1.05),  
3                 ylim=(0.001, 1.05))
```



The Area under the curve for the TPR-FPR curve is about 90.1%

In [85]:

```
1  ## adding the precision and recall for Logistic Regression to the comparison table
2  model_comparison_df.loc[[2], 'precision'] = metrics.precision_score(Y_test, prediction_lr)
3  model_comparison_df.loc[[2], 'recall'] = metrics.recall_score(Y_test, prediction_lr)
4  model_comparison_df
```

Out[85]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381
2	0.891156	0.108844	0.900605	Logistic Regression	0.708333	0.404762

## K - Nearest Neighbor

In [86]:

```
1  # creating a KNN classifier where K = 1
2  knn_1_model = KNeighborsClassifier(n_neighbors=1)
3  print("Deafult Parameters:\n\n" + str(knn_1_model.get_params(True)))
```

Deafult Parameters:

```
{'algorithm': 'auto', 'leaf_size': 30, 'metric': 'minkowski', 'metric_params': None, 'n_jobs': None, 'n_neighbors': 1, 'p': 2, 'weights': 'uniform'}
```

In [87]:



```
1 # fitting the KNN model with the training data and training output
2 knn_1_model.fit(X_train, Y_train)
3
4 # calculation of Training score
5 training_score_knn = knn_1_model.score(X_train,Y_train)
6 print("Training score :" + str(training_score_knn))
7
8 # calculation of Test score
9 testing_score_knn = knn_1_model.score(X_test,Y_test)
10 print("Testing score :" + str(testing_score_knn))
```

Training score :1.0

Testing score :0.891156462585034

The training score is much higher than the testing score indicating an overfit

### Choosing a optimal value for K

In [88]:




```
1 # Computing the cross validation accuracy for different values of K
2 accuracy_rate = []
3
4 for i in range(1,40):
5
6     knn = KNeighborsClassifier(n_neighbors=i)
7     score=cross_val_score(knn,X_train,Y_train,cv=10)
8     accuracy_rate.append(score.mean())
```

In [89]:

```
1 # Plotting the Accuracy Rate
2 plt.plot(range(1,40),accuracy_rate,color='blue', linestyle='dashed', marker='o',
3          markerfacecolor='red', markersize=10)
4 plt.title('Accuracy Rate vs. K Value')
5 plt.xlabel('K')
6 plt.ylabel('Accuracy Rate')
```

Out[89]:

Text(0, 0.5, 'Accuracy Rate')



In [90]:


```
1 # Computing the cross validation error rate (i.e 1-accuracy_rate) for different values
2 error_rate = []
3
4 for i in range(1,40):
5
6     knn = KNeighborsClassifier(n_neighbors=i)
7     score=cross_val_score(knn,X_train,Y_train,cv=10)
8     error_rate.append(1-score.mean())
```

In [91]:

```
1 # Plotting the error_rate
2
3 plt.plot(range(1,40),error_rate,color='blue', linestyle='dashed', marker='o',
4          markerfacecolor='red', markersize=10)
5 plt.title('Error Rate vs. K Value')
6 plt.xlabel('K')
7 plt.ylabel('Error Rate')
```

Out[91]:

Text(0, 0.5, 'Error Rate')



- From the plot of accuracy rate, we see that beyond K=20, the accuracy is varies slightly between 0.870 and 0.880
- From the plot of error rate, we see that beyonf K=20, the error rate varies 0.120 and 0.130
- so we choose K = 21

In [92]:



```

1 # Creating an instance of K-Nearest Neighbor with K= 21
2 knn_20_model = KNeighborsClassifier(n_neighbors=21)
3 print("Deafult Parameters:\n\n\n" + str(knn_20_model.get_params(True)))
4
5 # Adding a row for K-Nearest Neighbor in the model_comparison table
6 model_comparison_df = model_comparison_df.append({'model': 'K-Nearest Neighbor'}, ignore
7 model_comparison_df

```

Deafult Parameters:

```
{'algorithm': 'auto', 'leaf_size': 30, 'metric': 'minkowski', 'metric_params': None, 'n_jobs': None, 'n_neighbors': 21, 'p': 2, 'weights': 'uniform'}
```

Out[92]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381
2	0.891156	0.108844	0.900605	Logistic Regression	0.708333	0.404762
3	NaN	NaN	NaN	K-Nearest Neighbor	NaN	NaN

## Training the model and Prediction

In [93]:



```

1 # fitting the KNN model with the training data and training output
2 knn_20_model.fit(X_train, Y_train)
3
4 # predicting the values for X-test
5 prediction_knn= knn_20_model.predict(X_test)

```

## Evaluating the model

### a) Classification Report



In [94]:



```

1 # predicting the values for X-test
2 knn_report = classification_report(Y_test, prediction_knn,output_dict=True)
3
4 print(classification_report(Y_test, prediction_cart))

```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	252
1	0.76	0.31	0.44	42
accuracy			0.89	294
macro avg	0.83	0.65	0.69	294
weighted avg	0.88	0.89	0.87	294

**Accuracy of the KNN Classifier is 89%**

## b) Confusion Matrix


In [95]:



```

1 conf_mat = create_conf_mat(Y_test, prediction_knn)
2 sns.heatmap(conf_mat, annot=True, fmt='d', cbar=False)
3 plt.xlabel('Predicted Values')
4 plt.ylabel('Actual Values')
5 plt.title('Actual vs. Predicted Confusion Matrix')
6 plt.show()

```



## c) Accuracy

In [96]:



```

1 accuracy_knn = knn_20_model.score(X_test, Y_test)
2
3 print("Mean accuracy on the test set (when using KNN Classifier):\n {0:.3f}"\
4       .format(accuracy_knn))
5
6 ## adding the auc to the comparison table
7 model_comparison_df.loc[[3], 'accuracy'] = accuracy_knn
8 model_comparison_df

```

Mean accuracy on the test set (when using KNN Classifier):  
0.881

Out[96]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381
2	0.891156	0.108844	0.900605	Logistic Regression	0.708333	0.404762
3	0.880952	NaN	NaN	K-Nearest Neighbor	NaN	NaN

#### d) Error Rate

In [97]:



```

1 test_error_rate_knn = 1 - accuracy_knn
2 print("The test error rate for the KNN model is:\n {0: .4f}"\
3       .format(test_error_rate_knn))
4
5 ## adding the error_rate to the comparison table
6 model_comparison_df.loc[[3], 'error_rate'] = test_error_rate_knn
7 model_comparison_df

```

The test error rate for the KNN model is:  
0.1190

Out[97]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381
2	0.891156	0.108844	0.900605	Logistic Regression	0.708333	0.404762
3	0.880952	0.119048	NaN	K-Nearest Neighbor	NaN	NaN

#### e) Area under the Curve

In [98]:



```

1 predictions_prob_knn = knn_20_model.predict_proba(X_test)[:, 1]
2
3 # calculating the FPR and TPR
4 fpr_knn, tpr_knn, _knn = roc_curve(Y_test,
5                                   predictions_prob_knn,
6                                   pos_label = 1)
7
8 # calculating the area under the curve
9 auc_knn = auc(fpr_knn, tpr_knn)
10
11 ## adding the auc to the comparison table
12 model_comparison_df.loc[[3], 'auc'] = auc_knn
13 model_comparison_df

```

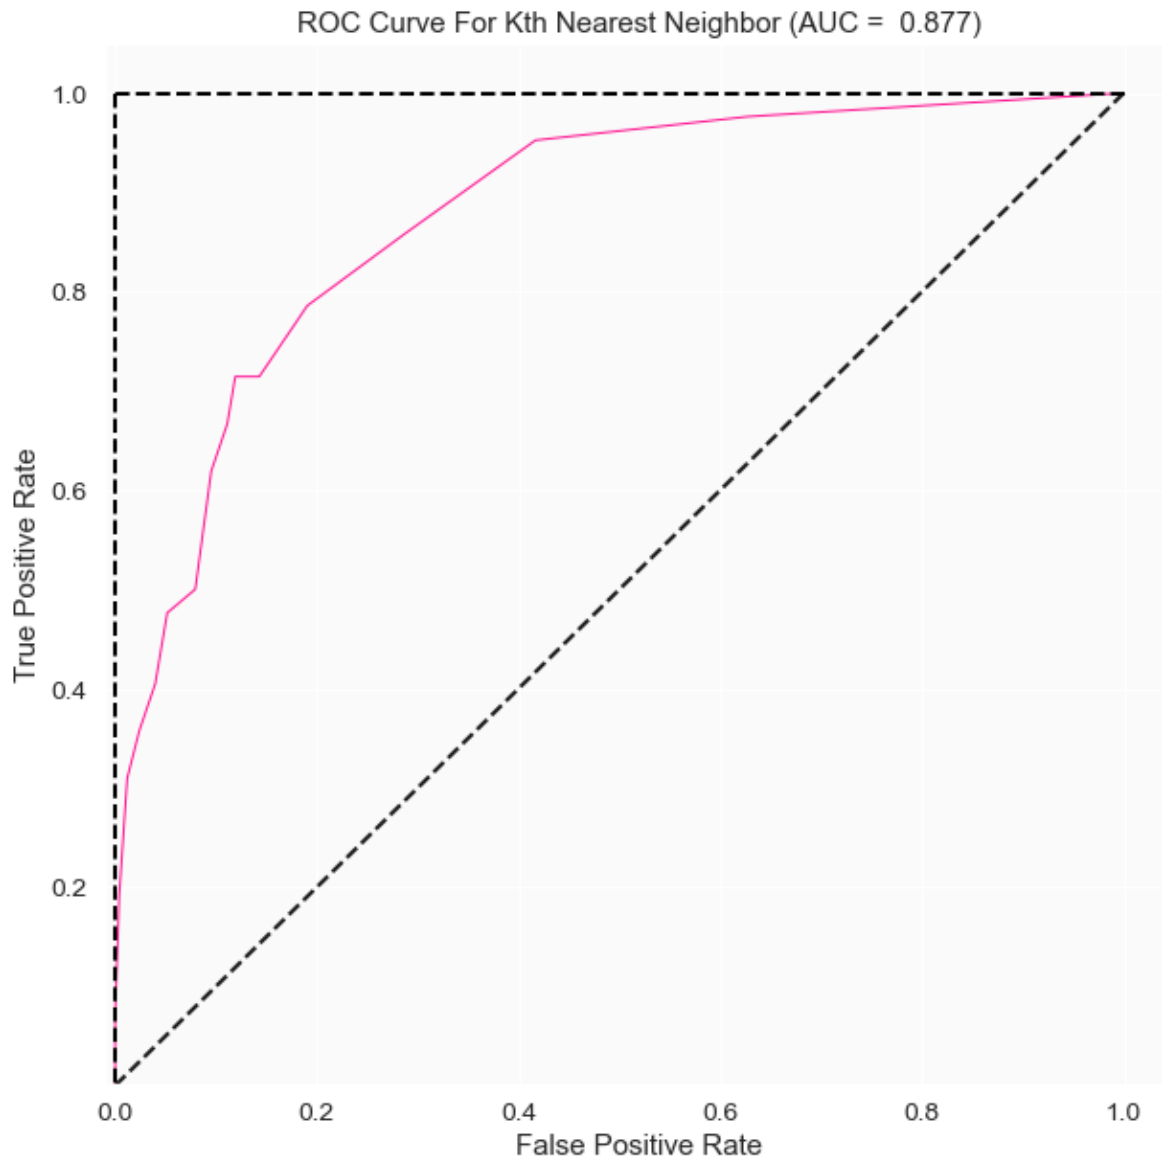
Out[98]:

	accuracy	error_rate	auc	model	precision	recall
0	0.887755	0.112245	0.849254	Decision Tree Classifier	0.764706	0.309524
1	0.904762	0.0952381	0.913076	Random Forest Classifier	0.791667	0.452381
2	0.891156	0.108844	0.900605	Logistic Regression	0.708333	0.404762
3	0.880952	0.119048	0.876701	K-Nearest Neighbor	NaN	NaN

In [99]:



```
1 plot_roc_curve(fpr_knn, tpr_knn, auc_knn, 'knn',  
2                 xlim=(-0.01, 1.05),  
3                 ylim=(0.001, 1.05))
```



The Area under the curve for the TPR-FPR curve is about 87.7%

In [100]:

```
1 ## adding the precision and recall for Logistic Regression to the comparison table
2 model_comparison_df.loc[[3], 'precision'] = metrics.precision_score(Y_test, prediction_kr
3 model_comparison_df.loc[[3], 'recall'] = metrics.recall_score(Y_test, prediction_knn)
```

## Comparison of Machine Learning Models

In [102]:

```
1 # setting model as the index
2 #model_comparison_df.drop(columns = ['Accuracy', 'Error_Rate', 'AUC'], inplace=True)
3 model_comparison_df.set_index(["model"], inplace = True)
4 HTML(model_comparison_df.to_html(classes='table-bordered'))
5
```

Out[102]:

	accuracy	error_rate	auc	precision	recall
model					
<b>Decision Tree Classifier</b>	0.887755	0.112245	0.849254	0.764706	0.309524
<b>Random Forest Classifier</b>	0.904762	0.0952381	0.913076	0.791667	0.452381
<b>Logistic Regression</b>	0.891156	0.108844	0.900605	0.708333	0.404762
<b>K-Nearest Neighbor</b>	0.880952	0.119048	0.876701	0.629630	0.404762

## Conclusion

- In terms of model accuracy, the best model for predicting the wine quality (as Good and Bad) is Random Forest Classifier
- The top three parameters which determine the quality of wine are Alcohol, Sulphates and Volatile Acidity

